


# MULTICORE PROGRAMMING WITH OPENMP



26.03.2024

**Philip Wiese**

wiesep@iis.ee.ethz.ch

**Zexin Fu**

zexifu@iis.ee.ethz.ch

Original slides by Bjoern Forsberg (bjoernf@iis.ee.ethz.ch)  
Based on slides by Dr. Andrea Marongiu (a.marongiu@unibo.it)

# Parts of exercise




- Understand OpenMP parallelization
  - ▣ Annotations, Compiler transformation, library support
  
- Exercises on core concepts, such as
  - ▣ SPMD parallelization (OpenMP parallel for)
    - Work distribution and scheduling
  - ▣ MPMD parallelization (OpenMP sections and tasks)
  - ▣ Data races
  
- Parallelization of simple Convolution kernel

# Programming model: OpenMP


- De-facto standard for the **shared memory** programming model
- A collection of **compiler directives**, **library routines** and **environment variables**
- Easy to specify parallel execution within a **serial code**
- Requires **special support** in the compiler
- Generates calls to **threading libraries** (e.g. pthreads)
- Focus on **loop-level** parallel execution (still the case?)
- Popular in high-end embedded

# Fork/Join Parallelism

## Sequential program



## Parallel program



- Initially only master thread is active
- Master thread executes sequential code
- Fork: Master thread creates or awakens additional threads to execute parallel code
- Join: At the end of parallel code created threads are suspended upon **barrier** synchronization

# Pragmas

- **Pragma**: a compiler directive in C or C++
- Stands for “pragmatic information”
- A way for the programmer to communicate with the compiler
- Compiler free to ignore pragmas: original sequential semantic is not altered
- Syntax:

**#pragma omp** *<rest of pragma>*

Example:

**#pragma omp** parallel for num\_threads(4)

How many  
threads

# Components of OpenMP

## a subset of the directives

### *Directives*

- ❖ Parallel regions
  - `#pragma omp parallel`
- ❖ Work sharing
  - `#pragma omp for`
  - `#pragma omp sections`
- ❖ Synchronization
  - `#pragma omp barrier`
  - `#pragma omp critical`
  - `#pragma omp atomic`

### *Runtime Library*

### *Clauses*

- ❖ Data scope attributes
  - `private`
  - `shared`
  - `reduction`
- ❖ Loop scheduling
  - `static`
  - `dynamic`

- ❖ Thread Forking/Joining
  - `omp_parallel_start()`
  - `omp_parallel_end()`
- ❖ Loop scheduling
- ❖ Thread IDs
  - `omp_get_thread_num()`
  - `omp_get_num_threads()`

# Outlining parallelism

## The `parallel` directive

- Fundamental construct to outline parallel computation within a sequential program
- Code within its scope is **replicated** among threads
- Defers implementation of parallel execution to the runtime (machine-specific, e.g. `pthread_create`)

**A sequential program..  
..is easily parallelized**

```
int main()  
{  
  #pragma omp parallel  
  {  
    printf ("\nHello world!");  
  }  
}
```

# Outlining parallelism

## The `parallel` directive

- Fundamental construct to outline parallel computation within a sequential program
- Code within its scope is **replicated** among threads
- Defers implementation of parallel execution to the runtime (machine-specific, e.g. `pthread_create`)

**A sequential program..  
..is easily parallelized**

```
int main()
{
    #pragma omp parallel
    {
        printf ("\nHello world!");
    }
}
```

```
int main()
{
    omp_parallel_start(&parfun, ...);
    parfun();
    omp_parallel_end();
}

int parfun(...)
{
    printf ("\nHello world!");
}
```



# #pragma omp parallel

Code originally contained within the scope of the pragma is outlined to a new function within the compiler

```
int main()
{
    #pragma omp parallel
```

```
    printf ("\nHello world!");
}
```

```
int main()
{
    omp_parallel_start(&parfun, ...);
    parfun();
    omp_parallel_end();
}
```

```
int parfun(...)
{
    printf ("\nHello world!");
}
```

# #pragma omp parallel

The **#pragma** construct in the **main** function is replaced with function calls to the runtime library

```
int main()
{
#pragma omp parallel
{
    printf ("\nHello world!");
}
}
```

```
int main()
{
    omp_parallel_start(&parfun, ...);
    parfun();
    omp_parallel_end();
}

int parfun(...)
{
    printf ("\nHello world!");
}
```

# #pragma omp parallel

First we call the runtime to fork new threads, and pass them a pointer to the function to execute in parallel

```
int main()
{
    #pragma omp parallel
    {
        printf ("\nHello world!");
    }
}
```

```
int main()
{
    omp_parallel_start(&parfun, ...);
    parfun();
    omp_parallel_end();
}

int parfun(...)
{
    printf ("\nHello world!");
}
```

# #pragma omp parallel

Then the master itself calls  
the parallel function




```
int main()
{
    #pragma omp parallel
    {
        printf ("\nHello world!");
    }
}
```

```
int main()
{
    omp_parallel_start(&parfun, ...);
    parfun();
    omp_parallel_end();
}

int parfun(...)
{
    printf ("\nHello world!");
}
```

# #pragma omp parallel

Finally we call the runtime to synchronize threads with a barrier and suspend them



```
int main()
{
    #pragma omp parallel
    {
        printf ("\nHello world!");
    }
}
```

```
int main()
{
    omp_parallel_start(&parfun, ...);
    parfun();
    omp_parallel_end();
}

int parfun(...)
{
    printf ("\nHello world!");
}
```

# #pragma omp parallel

## Data scope attributes

```
int main()
```

```
{
```

```
    int id;
```

```
    int a = 5;
```

```
    #pragma omp parallel
```

```
    {
```

```
        id = omp_get_thread_num();
```

```
        if (id == 0)
```

```
            printf ("Master: a = %d.", a*2);
```

```
        else
```

```
            printf ("Slave: a = %d.", a);
```

```
    }
```

```
}
```

Call runtime to get thread ID:  
Every thread sees a different value

Master and slave threads  
access the same variable **a**

A slightly more complex example

# #pragma omp parallel

## Data scope attributes

```
int main()
```

```
{
```

```
    int id;
```

```
    int a = 5;
```

```
    #pragma omp parallel
```

```
    {
```

```
        id = omp_get_thread_num();
```

```
        if (id == 0)
```

```
            printf ("Master: a = %d.", a);
```

```
        else
```

```
            printf ("Slave: a = %d.", a);
```

```
    }
```

```
}
```

Call runtime to get thread ID:

Every thread sees a different value

How to inform the compiler  
about these different  
behaviors?

Master and slave threads  
access the same variable **a**

A slightly more complex example

# #pragma omp parallel

## Data scope attributes

What is the view of memory among different threads in a parallel region?

```
int main()
{
    int id;
    int a = 5;
    #pragma omp parallel shared (a) private (id)
    {
        id = omp_get_thread_num();
        if (id == 0)
            printf ("Master: a = %d.", a*2);
        else
            printf ("Slave: a = %d.", a);
    }
}
```

Insert code to retrieve the address  
of the shared object from within  
each parallel thread

Allow symbol privatization:  
Each thread contains a  
private copy of this variable

A slightly more complex example



# #pragma omp parallel

## Data scope attributes

```
int main()
{
    int id;
    int a = 1;
    #pragma omp parallel
    {
        id = omp_get_thread_num();
        if (id == 0)
            printf("Thread 0\n");
        else
            printf("Thread %d\n", id);
    }
}
```

### Correctness issues

What if:

- **a** was not marked for shared access?
- **id** was not marked for private access?

Insert code to retrieve the address of the shared object from within a parallel thread

(id)

Allow symbol privatization: each thread contains a private copy of this variable

A slightly more complex example

# More data sharing clauses

---

- firstprivate
  - ▣ copyin, private storage
- lastprivate
  - ▣ copyout, private storage

# SPMD VS MPMD

Recall..

- SPMD (single program, multiple data)
  - ▣ Processors execute the same stream of instructions over different data
  - ▣ `#pragma omp for`
- MPMD (multiple program, multiple data)
  - ▣ Processors execute different streams of instructions over (possibly) different data
  - ▣ `#pragma omp sections`
  - ▣ `#pragma omp task`

# Sharing work among threads

## The **for** directive

---

- The **parallel** pragma instructs every thread to execute all of the code inside the block
- If we encounter a **for** loop that we want to divide among threads, we use the **for** pragma

**#pragma omp for**

# #pragma omp for

**#pragma omp for** can be placed everywhere inside a **parallel** construct, or combined with it, as in the example

The code of the **for** loop is moved inside the outlined function.

```
int main()
{
    #pragma omp parallel for
    {
        for (i=0; i<10; i++)
            a[i] = i;
    }
}
```

```
int main()
{
    omp_parallel_start(&parfun, ...);
    parfun();
    omp_parallel_end();
}

int parfun(...)
{
    int LB = ...;
    int UB = ...;

    for (i=LB; i<UB; i++)
        a[i] = i;
}
```

# The `schedule` clause

## Static Loop Partitioning

```
#pragma omp for schedule(static)
{
    for (i=0; i<12; i++)
        a[i] = i;
}
```

Useful for:

- *Simple, regular loops*
- *Iterations with equal duration*

Es. 12 iterations (N), 4 threads (Nthr)

**DATA CHUNK**

$$C = \text{ceil} \left( \frac{N}{N_{\text{thr}}} \right)$$

**3  
iterations  
thread**

Iteration space



**LOWER BOUND**

$$LB = C * TID$$

**UPPER BOUND**

$$UB = \min \{ [C * (TID + 1)], N \}$$

Thread ID (TID)

0

1

2

3

0

3

6

9

3

6

9

12

# The schedule clause

## Static Loop Partitioning

Es. 12 iterations (N), 4 threads (Nthr)

```
#pragma omp for schedule(static)
{
    for (i=0; i<12; i++)
        a[i] = i;
}
```

Useful for:

- Simple, regular loops
- Iterations with equal duration

**DATA CHUNK**

$$\text{chunk} \left( \frac{N}{N_{\text{thr}}} \right)$$

**3  
iterations  
thread**

**What happens with  
static scheduling  
when iterations have  
different duration?**

**LOWER BOUND**

**UPPER BOUND**

$\{ \lfloor \frac{N}{N_{\text{thr}}} \cdot (TID + 1) \rfloor, N \}$

	0	1	2	3
	0	3	6	9
	3	6	9	12

# The `schedule` clause

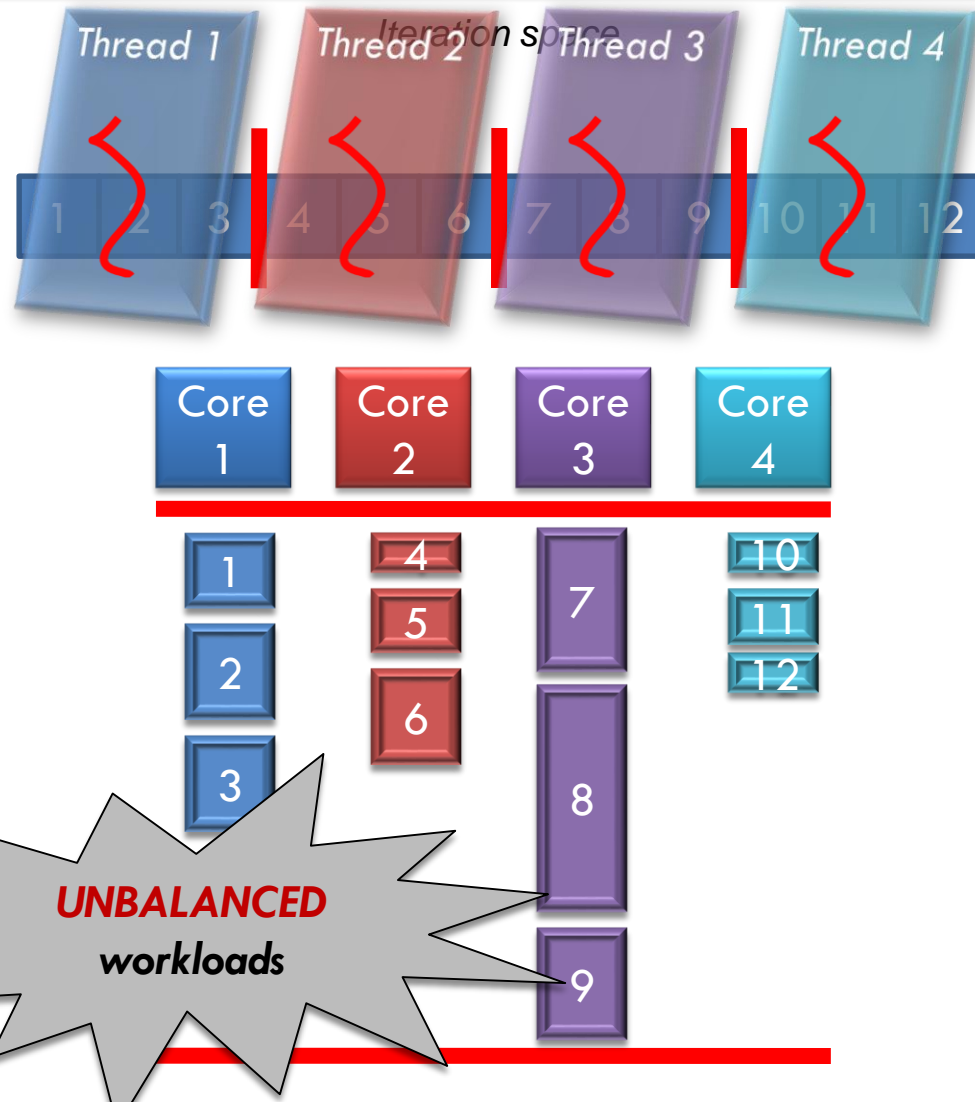
## Static Loop Partitioning

```
#pragma omp for schedule(static)
{
    for (i=0; i<12; i++)
    {
        int start = rand();
        int count = 0;

        while (start++ < 256)
            count++;

        a[count] = foo();
    }
}
```

A variable amount of work in  
each iteration






# The `schedule` clause

## Dynamic Loop Partitioning

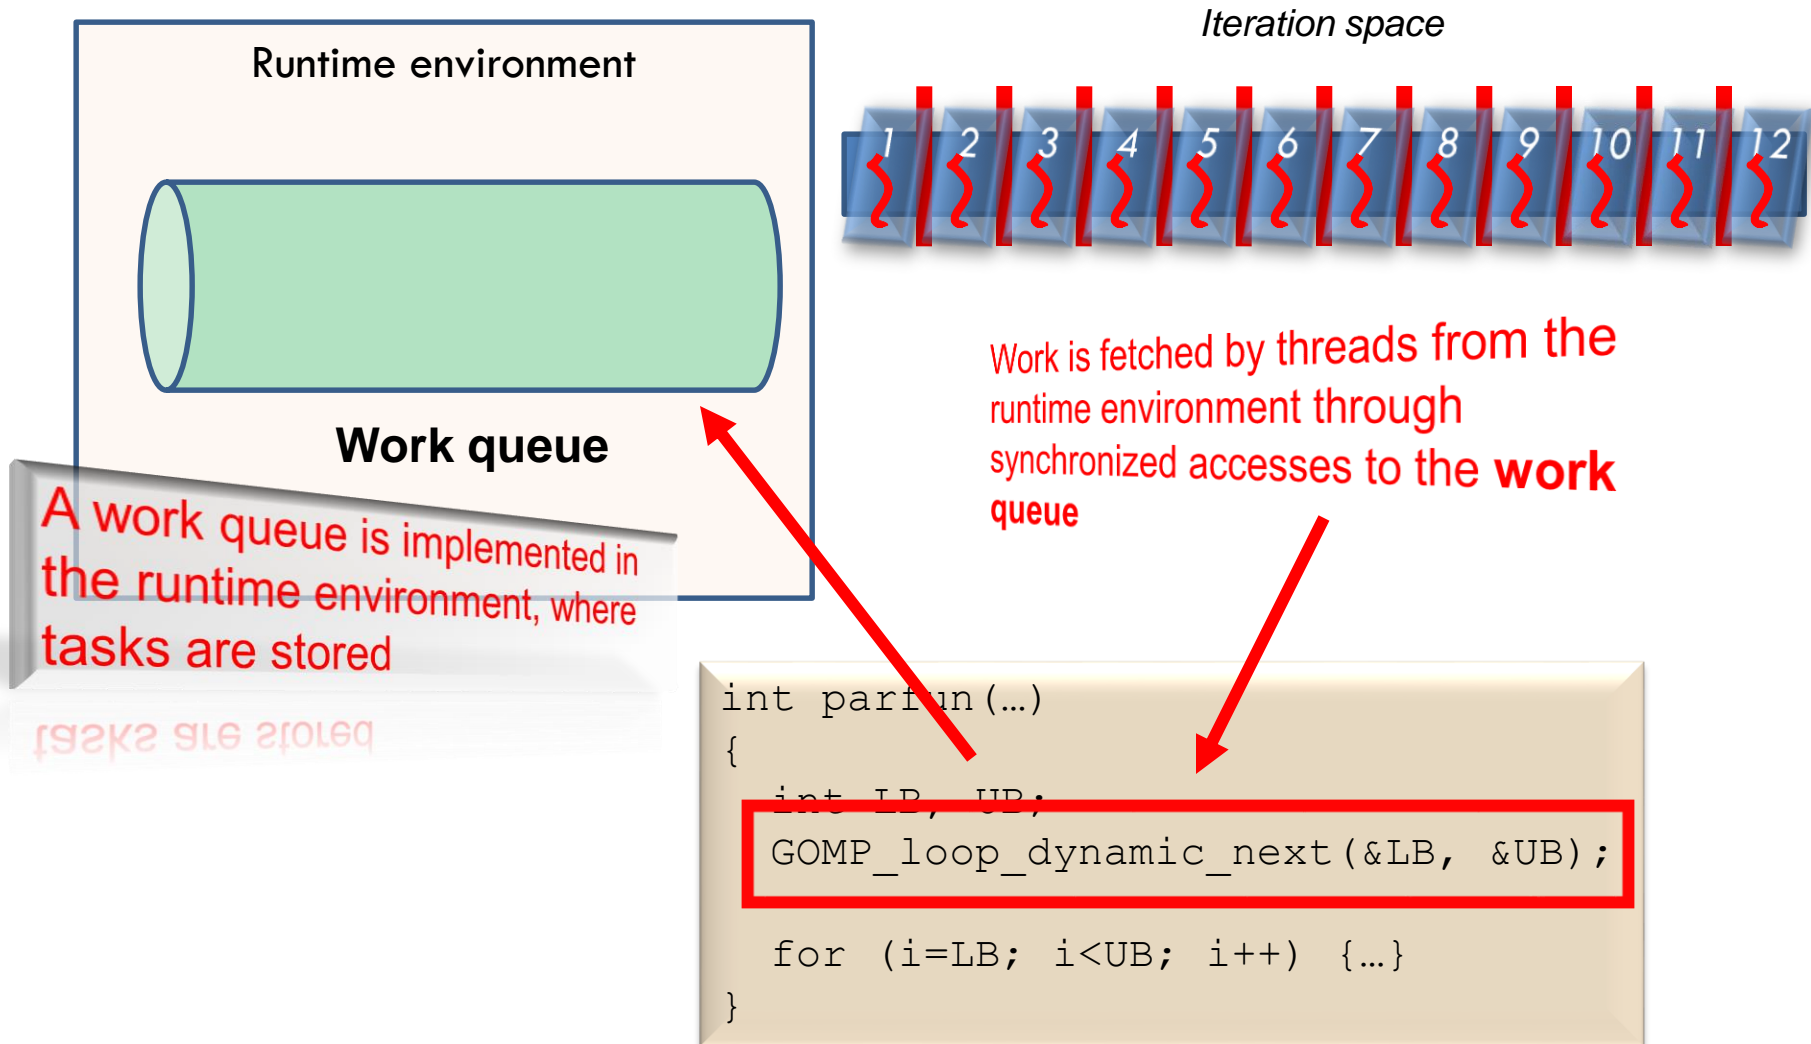
```
#pragma omp for schedule(dynamic  
{  
  for (i=0; i<12; i++)  
  {  
    int start = rand();  
    int count = 0;  
  
    while (start++ < 256)  
      count++;  
  
    a[count] = foo();  
  }  
}
```



A thread is generated for  
every single iteration

# The `schedule` clause

## Dynamic Loop Partitioning





# The `schedule` clause

## Dynamic Loop Partitioning

Remember results with static scheduling..

*Iteration space*



**Speedup!**

**BALANCED**  
workloads

# Parallelization granularity

## Iteration chunking

- Fine-grain Parallelism

- Best opportunities for load balancing, but..
- Small amounts of computational work between parallelism computation stages
- Low computation to parallelization ratio → High parallelization overhead



- Coarse-grain Parallelism


- Harder to load balance efficiently, but..
- Large amounts of computational work between parallelism computation stages
- High computation to parallelization ratio → Low parallelization overhead



# The `schedule` clause

## Dynamic Loop Partitioning


Remember results with  
static scheduling..



*Iteration space*



```
#pragma omp for schedule(dynamic, 1)
```



# The `schedule` clause



## Dynamic Loop Partitioning

chunking overhead

*Iteration space*




```
#pragma omp for schedule(dynamic, 1)
```



# The `schedule` clause

## Dynamic Loop Partitioning



# #pragma omp barrier

- Most important **synchronization** mechanism in shared memory fork/join parallel programming
- All threads participating in a parallel region wait until everybody has finished before computation flows on
- This prevents later stages of the program to work with **inconsistent** shared data
- It is implied at the end of **parallel** constructs, as well as **for** and **sections** (unless a **nowait** clause is specified)




# #pragma omp critical

- **Critical Section:** a portion of code that only one thread at a time may execute
- We denote a critical section by putting the pragma

**#pragma omp critical**

in front of a block of C code

# $\pi$ -finding code example




```
double area, pi, x;
int i, n;
#pragma omp parallel for private(x) \
                        shared(area)
{
    for (i=0; i<n; i++) {
        x = (i + 0.5)/n;
        area += 4.0/(1.0 + x*x);
    }
    pi = area/n;
}
```

Synchronize accesses to  
shared variable **area** to  
avoid inconsistent results

# Race condition (Cont'd)


- Thread A reads “11.667” into a local register
- Thread B reads “11.667” into a local register
- Thread A updates area with “11.667+3.765”
- Thread B ignores write from thread A and updates area with “11.667 + 3.563”

time



# $\pi$ -finding code example

```
double area, pi, x;
int i, n;
#pragma omp parallel for private(x) shared(area)
{
    for (i=0; i<n; i++) {
        x = (i + 0.5)/n;
        #pragma omp critical
        area += 4.0/(1.0 + x*x);
    }
}
pi = area/n;
```



**#pragma omp critical** protects the code within its scope by acquiring a lock before entering the critical section and releasing it after execution

# $\pi$ -finding code example


```
double area, pi, x;
int i, n;
#pragma omp parallel for \
    private(x) \
    shared(area)
{
    for (i=0; i<n; i++) {
        x = (i + 0.5) / n;
        #pragma omp critical
        area += 4.0 / (1.0 + x*x);
    }
    pi = area / n;
}
```

**Parallel**

**Sequential**



Waiting for lock



# Correctness, not performance!

- A programming pattern such as `area += 4.0 / (1.0 + x*x) ;` in which we:
  - ▣ *Fetch the value of an operand*
  - ▣ *Add a value to it*
  - ▣ *Store the updated value*is called a **reduction**, and is commonly supported by parallel programming APIs
- OpenMP takes care of storing partial results in **private variables** and combining partial results after the loop

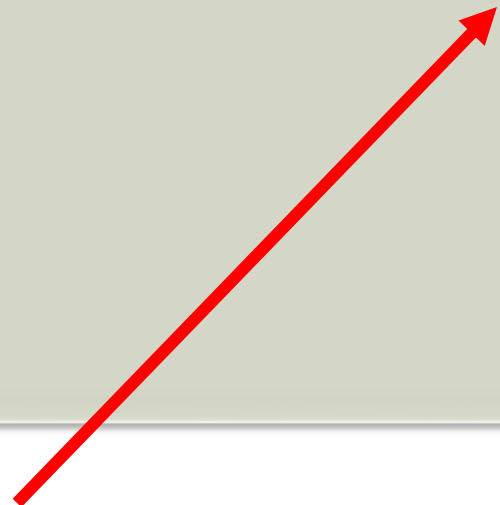
# Correctness, not performance!

- As a matter of fact, using locks makes execution **sequential**
- To dim this effect we should try use **fine grained locking** (i.e. make critical sections as small as possible)
- A simple instruction to compute the value of area in the previous example is translated into many more simpler instructions within the compiler!
- The programmer is not aware of the real **granularity** of the critical section

# Correctness, not performance!

```
double area, pi, x;
int i, n;
#pragma omp parallel for private(x) shared(area) reduction(+:area)
{
    for (i=0; i<n; i++) {
        x = (i + 0.5)/n;


        area += 4.0/(1.0 + x*x);
    }
}
pi = area/n;
```





The **reduction** clause instructs the compiler to create **private** copies of the **area** variable for every thread. At the end of the loop partial sums are combined on the shared **area** variable



# Correctness, not performance!



reduction (+: area)



# SPMD VS MPMD

Recall..

- SPMD (single program, multiple data)
  - ▣ Processors execute the same stream of instructions over different data
  - ▣ `#pragma omp for`
- MPMD (multiple program, multiple data)
  - ▣ Processors execute different streams of instructions over (possibly) different data
  - ▣ `#pragma omp sections`
  - ▣ `#pragma omp task`

# Sharing work among threads

## The `sections` directive

- The `for` pragma allows to exploit **data parallelism** in loops
- OpenMP also provides directives to exploit **task parallelism**

`#pragma omp sections`

# Task Parallelism Example

```
int main()
{

    v = alpha();

    w = beta ();


    y = delta ();

    x = gamma (v, w);

    z = epsilon (x, y));

    printf ("%f\n", z);
}
```

Identify independent nodes  
in the task graph, and outline  
parallel computation with the  
**sections** directive



# Task Parallelism Example

```
int main()
{
#pragma omp parallel sections {

    v = alpha();

    w = beta ();

}
#pragma omp parallel sections {

    y = delta ();


    x = gamma (v, w);

}
    z = epsilon (x, y));

printf ("%f\n", z);
}
```

Barriers implied here!

FIRST SOLUTION




# Task Parallelism Example

```
int main()
{
    #pragma omp parallel sections {
        #pragma omp section
        v = alpha();
        #pragma omp section
        w = beta ();
    }
    #pragma omp parallel sections {
        #pragma omp section
        y = delta ();
        #pragma omp section
        x = gamma (v, w);
    }
    z = epsilon (x, y);

    printf ("%f\n", z);
}
```

Each parallel task within a  
sections block identifies a  
section



# Task Parallelism Example

```
int main()
{
    v = alpha();

    w = beta ();

    y = delta ();


    x = gamma (v, w);

    z = epsilon (x, y));

    printf ("%f\n", z);
}
```

**SECOND  
SOLUTION**


Identify independent nodes  
in the task graph, and outline  
parallel computation with the  
**sections** directive



# Task Parallelism Example

```
int main()  
{  
    #pragma omp parallel sections {  
  
        v = alpha();  
  
        w = beta ();  
  
        y = delta ();  
  
    }  
  
    x = gamma (v, w);  
  
    z = epsilon (x, y));  
  
    printf ("%f\n", z);  
}
```

Each parallel task within a  
**sections** block identifies a  
**section**





# Task Parallelism Example


```
int main()
{
    #pragma omp parallel sections {
        #pragma omp section
        v = alpha();
        #pragma omp section
        w = beta ();
        #pragma omp section
        y = delta ();

        x = gamma (v, w);

        z = epsilon (x, y));

    printf ("%f\n", z);
}
```

Each parallel task within a  
sections block identifies a  
section



# Task parallelism

- The **sections** directive allows a very limited form of task parallelism
- All tasks must be statically outlined in the code
  - ▣ What if a functional loop (**while**) body is identified as a task?
    - **Unrolling?** Not feasible for high iteration count
  - ▣ What if recursion is used?

# What is an OpenMP task?

- Tasks are work units which execution **may** be deferred
  - ▣ they can also be executed immediately!
- Tasks are composed of:
  - ▣ code to execute
  - ▣ data environment
    - Initialized at creation time
  - ▣ internal control variables (ICVs)

# Task directive

```
#pragma omp task [ clauses ]  
    structured block
```

- Each encountering thread creates a task
  - ▣ Packages code and data environment
- Highly composable. Can be nested
  - ▣ inside parallel regions
  - ▣ inside other tasks
  - ▣ inside worksharing constructs (for, sections)

# List traversal with tasks

- Why?
  - ▣ Example: **list traversal**

## EXAMPLE

```
void traverse_list (List l)
{
    Element e ;

    for ( e = e→first; e; e = e→next )
    #pragma omp task
        process ( e ) ;
}
```

What is the scope of e?



# Task data scoping

- Data scoping clauses
  - ▣ `shared(list)`
  - ▣ `private(list)`
  - ▣ `firstprivate(list)`
    - data is captured at creation
  - ▣ `default(shared | none)`

# Task data scoping

when there are no clauses..

- If no clause
  - ▣ Implicit rules apply
    - e.g., global variables are shared
- Otherwise...
  - ▣ firstprivate
  - ▣ shared attribute is lexically inherited

Tip

**default(none)** is your friend

Use it if you do not see it clear

# List traversal with tasks

## EXAMPLE

```
void traverse_list (List l)
{
    Element e ;

    for ( e = e→first; e; e = e→next )
    #pragma omp task
        process ( e ) ;
}
```

← **e is firstprivate**




# List traversal with tasks

## EXAMPLE

```
void traverse_list (List l)
{
    Element e ;

    for ( e = e→first; e; e = e→next )
    #pragma omp task
        process ( e ) ;
}
```



how we can guarantee here that the traversal is finished?

# Task synchronization

- Barriers (implicit or explicit)
  - ▣ All tasks created by any thread of the current team are guaranteed to be completed at barrier exit
- Task barrier
  - #pragma omp taskwait**
  - ▣ Encountering task suspends until **child** tasks complete
    - Only direct **childs**, not descendants!

# List traversal with tasks

## EXAMPLE

```
void traverse_list (List l)
{
    Element e ;

    for ( e = e→first; e; e = e→next )
    #pragma omp task
        process ( e ) ;

    #pragma omp taskwait
}
```



All tasks guaranteed to be completed here

# Task parallelism

## □ Why?

### ▣ Example: list traversal

**CAREFUL!**

- Multiple traversal of the same list

## EXAMPLE

```
List l;
```

```
#pragma omp parallel
```

```
traverse_list (l);
```

```
void traverse_list (List l)
{
    Element e ;

    for ( e = e→first; e; e = e→next )
#pragma omp task
        process ( e ) ;
}
```

# Task parallelism

## □ Why?

### ▣ Example: list traversal

#### EXAMPLE

```
List l;
```

```
#pragma omp parallel  
#pragma omp single  
traverse_list (l);
```

## Single traversal

- One thread enters single and creates all tasks
- All the team cooperates executing them

```
void traverse_list (List l)  
{  
    Element e ;  
  
    for ( e = e→first; e; e = e→next )  
#pragma omp task  
        process ( e ) ;  
}
```

# Task parallelism

- In case **task** is within a regular **counted loop** an alternative is to parallelize task creation among threads

## EXAMPLE

```
/* A DIFFERENT EXAMPLE */
```

```
#pragma omp parallel
```

```
Myfunc ();
```

## Multiple traversals

- Multiple threads create tasks
- All the team cooperates executing them

```
void Myfunc ()  
{  
    int i;  
    #pragma omp for  
        for (i=LB; i<UB; i++)  
    #pragma omp task  
        process ( i ) ;  
}
```



**EXERCISE!**